

semantic expressiveness increases →

Proxy metrics

fidelity, stability,
sparsity, cost
EXP2/EXP3
(this paper)

**Benchmark-
grounded**

synthetic or domain
ground truth

Expert judgment

plausibility,
rubric-based review

User-centered

task outcomes,
trust calibration

LLM-based semantic evaluation

more expressive than proxies, more scalable
than human review, but not self-validating

EXP4 (this paper): $\text{ICC}(1,1) < 0.75$ on
all seven dimensions, two judge panels

← scalability and reproducibility increase